

## Question 1

- Add a batch normalization layer after the first fully-connected layer(fc1) (8 points).
- Save the model after training(Checkout our tutorial on how to save your model). Be careful that batch normalization layer performs differently between training and evaluation process, make sure you understand how to convert your model between training mode and evaluation mode(you can find hints in my code).
- Observe the difference of final training/testing accuracy with/without batch normalization layer.

### Code:

## Q1


```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.conv2 = nn.Conv2d(32, 32, 3, padding=1)
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
        self.conv4 = nn.Conv2d(64, 64, 3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(64 * 8 * 8, 512)
        self.fc2 = nn.Linear(512, 10)
        self.bnorm = nn.BatchNorm1d(512)
```


```
    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.pool(x)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = self.bnorm(x)
        x = self.fc2(x)
        return x
```

```
    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```


**Result:**



Accuracy



Loss



**Question 2**

- Modify our model by adding another fully connected layer with 512 nodes at the second-to-last layer (before the fc2 layer) (8 points).
- Apply the model weights you saved at step 1 to initialize to the new model(only up to fc2 layer since after that all layers are newly created) before training.
- Train and save the model (Hint: check the end of the assignment description to see how to partially restore weights from a pretrained weights file).

**Code:**

## Q2-1

class Net(nn.Module):

def \_\_init\_\_(self):

super(Net, self).\_\_init\_\_()

self.conv1 = nn.Conv2d(3, 32, 3, padding=1)

self.conv2 = nn.Conv2d(32, 32, 3, padding=1)

self.conv3 = nn.Conv2d(32, 64, 3, padding=1)

self.conv4 = nn.Conv2d(64, 64, 3, padding=1)

self.pool = nn.MaxPool2d(2, 2)

self.fc1 = nn.Linear(64 \* 8 \* 8, 512)

self.fc\_q2 = nn.Linear(512, 512)

self.fc2 = nn.Linear(512, 10)

self.bnorm = nn.BatchNorm1d(512)

def forward(self, x):

x = F.relu(self.conv1(x))

x = F.relu(self.conv2(x))

x = self.pool(x)

x = F.relu(self.conv3(x))

x = F.relu(self.conv4(x))

x = self.pool(x)

x = x.view(-1, self.num\_flat\_features(x))

x = F.relu(self.fc1(x))

x = self.bnorm(x)

x = F.relu(self.fc\_q2(x))

x = self.fc2(x)

return x

def num\_flat\_features(self, x):

size = x.size()[1:] # all dimensions except the batch dimension

num\_features = 1

for s in size:

num\_features \*= s

return num\_features

## Q2-2

def partially\_restore\_weights(filepath):

pretrained\_dict = torch.load(filepath)

model\_dict = net.state\_dict()


pretrained\_dict = {key: val for key, val in pretrained\_dict.items() if key in model\_dict}

model\_dict.update(pretrained\_dict)


net.load\_state\_dict(model\_dict)

return net


**Result:**



## Accuracy



## Loss




### Question 3

- Try to use an adaptive schedule to tune the learning rate, you can choose from RMSprop, Adagrad and Adam (Hint: you don't need to implement any of these, look at Pytorch documentation please) (8 points).


#### Code:

```
criterion = nn.CrossEntropyLoss()
# optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.9)
optimizer = optim.Adam(net.parameters(), lr=0.001, amsgrad=True) # Q3 & Q4
```


#### Result:



#### Accuracy



#### Loss



**Question 4**

- Try to tune your network in another way (e.g. add/remove a layer, change the activation function, add/remove regularizer, change the number of hidden units, more batch normalization layers) not described in the previous four. You can start from random initialization or previous results as you wish (8 points).

**Code:**

## Q4

class Net(nn.Module):

def \_\_init\_\_(self):

super(Net, self).\_\_init\_\_()

self.conv1 = nn.Conv2d(3, 32, 3, padding=1)

self.conv2 = nn.Conv2d(32, 32, 3, padding=1)

self.conv3 = nn.Conv2d(32, 64, 3, padding=1)

self.conv4 = nn.Conv2d(64, 64, 3, padding=1)

self.conv5 = nn.Conv2d(64, 128, 3, padding=1)

self.conv6 = nn.Conv2d(128, 128, 3, padding=1)

self.pool = nn.MaxPool2d(2, 2)

self.fc1 = nn.Linear(128 \* 4 \* 4, 512)

self.fc\_q2 = nn.Linear(512, 512)

self.fc2 = nn.Linear(512, 10)

self.bnorm1d1 = nn.BatchNorm1d(512)

self.bnorm1d2 = nn.BatchNorm1d(512)

self.bnorm2d1 = nn.BatchNorm2d(32)

self.bnorm2d2 = nn.BatchNorm2d(64)

self.bnorm2d3 = nn.BatchNorm2d(128)

self.dropout = nn.Dropout(0.25)

def forward(self, x):

# Understanding the Disharmony between Dropout and Batch Normalization by Variance Shift, <https://arxiv.org/abs/1801.05134>

x = F.relu(self.conv1(x))

x = self.bnorm2d1(x)

x = F.relu(self.conv2(x))

x = self.pool(x)

x = self.dropout(x)

x = F.relu(self.conv3(x))

x = self.bnorm2d2(x)

x = F.relu(self.conv4(x))

x = self.pool(x)

x = self.dropout(x)

x = F.relu(self.conv5(x))

x = self.bnorm2d3(x)

x = F.relu(self.conv6(x))

x = self.pool(x)

x = x.view(-1, self.num\_flat\_features(x))

x = F.relu(self.fc1(x))

x = self.bnorm1d1(x)

x = F.relu(self.fc\_q2(x))

x = self.bnorm1d2(x)

x = self.fc2(x)

return x

def num\_flat\_features(self, x):

size = x.size()[1:] # all dimensions except the batch dimension



## CS535 DL Assignment3 Dongkyu Kim

```
num_features = 1
for s in size:
    num_features *= s
return num_features
```

# data augmentation for Q4

```
transform = transforms.Compose([
    transforms.RandomHorizontalFlip(), # randomly flip and rotate
    transforms.RandomRotation(10),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
```

```
criterion = nn.CrossEntropyLoss()
```


```
# optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.9)
```

```
optimizer = optim.Adam(net.parameters(), lr=0.001, amsgrad=True) # Q3 & Q4
```


### Result:



### Accuracy



### Loss



## ## Source Code

```
from __future__ import print_function
from __future__ import division
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
import torch.optim as optim
# from tensorboardX import SummaryWriter # for pytorch below 1.14
from torch.utils.tensorboard import SummaryWriter # for pytorch above or equal 1.14

# check if CUDA is available
train_on_gpu = torch.cuda.is_available()

if not train_on_gpu:
    print('CUDA is not available. Training on CPU ...')
else:
    print('CUDA is available! Training on GPU ...')

## Q4
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.conv2 = nn.Conv2d(32, 32, 3, padding=1)
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
        self.conv4 = nn.Conv2d(64, 64, 3, padding=1)
        self.conv5 = nn.Conv2d(64, 128, 3, padding=1)
        self.conv6 = nn.Conv2d(128, 128, 3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(128 * 4 * 4, 512)
        self.fc_q2 = nn.Linear(512, 512)
        self.fc2 = nn.Linear(512, 10)
        self.bnorm1d1 = nn.BatchNorm1d(512)
        self.bnorm1d2 = nn.BatchNorm1d(512)
        self.bnorm2d1 = nn.BatchNorm2d(32)
        self.bnorm2d2 = nn.BatchNorm2d(64)
        self.bnorm2d3 = nn.BatchNorm2d(128)
        self.dropout = nn.Dropout(0.25)

    def forward(self, x):
        # Understanding the Disharmony between Dropout and Batch Normalization by Variance Shift,
        # https://arxiv.org/abs/1801.05134
        x = F.relu(self.conv1(x))
        x = self.bnorm2d1(x)
```



### CS535 DL Assignment3 Dongkyu Kim

```
x = F.relu(self.conv2(x))
x = self.pool(x)
x = self.dropout(x)
x = F.relu(self.conv3(x))
x = self.bnorm2d2(x)
x = F.relu(self.conv4(x))
x = self.pool(x)
x = self.dropout(x)
x = F.relu(self.conv5(x))
x = self.bnorm2d3(x)
x = F.relu(self.conv6(x))
x = self.pool(x)
x = x.view(-1, self.num_flat_features(x))
x = F.relu(self.fc1(x))
x = self.bnorm1d1(x)
x = F.relu(self.fc_q2(x))
x = self.bnorm1d2(x)
x = self.fc2(x)
```

```
return x
```

```
def num_flat_features(self, x):
    size = x.size()[1:] # all dimensions except the batch dimension
    num_features = 1
    for s in size:
        num_features *= s
    return num_features
```

```
def eval_net(dataloader):
    correct = 0
    total = 0
    total_loss = 0
    net.eval() # Why would I do this?
    criterion = nn.CrossEntropyLoss(reduction='mean')
    for data in dataloader:
        images, labels = data
        images, labels = Variable(images).cuda(), Variable(labels).cuda()
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels.data).sum()
        loss = criterion(outputs, labels)
        total_loss += loss.item()
    net.train() # Why would I do this?
    return total_loss / total, correct.float() / total
```

```
if __name__ == "__main__":
    BATCH_SIZE = 32 #mini_batch size
```

### CS535 DL Assignment3 Dongkyu Kim

MAX\_EPOCH = 10 #maximum epoch to train

```
# transform = transforms.Compose(  
#     [transforms.ToTensor(),  
#     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))]) #torchvision.transforms.Normalize(mean, std)
```

```
# data augmentation for Q4  
transform = transforms.Compose([  
    transforms.RandomHorizontalFlip(), # randomly flip and rotate  
    transforms.RandomRotation(10),  
    transforms.ToTensor(),  
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))  
])
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,  
                                         download=True, transform=transform)  
trainloader = torch.utils.data.DataLoader(trainset, batch_size=BATCH_SIZE,  
                                           shuffle=True, num_workers=2)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False,  
                                         download=True, transform=transform)  
testloader = torch.utils.data.DataLoader(testset, batch_size=BATCH_SIZE,  
                                          shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat',  
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

```
print('Building model...')  
net = Net().cuda()  
# net = partially_restore_weights('mytraining1.pth') #Q2  
net.train() # Why would I do this?
```

```
images, labels = next(iter(trainloader))  
grid = torchvision.utils.make_grid([images.cuda()])
```

```
writer = SummaryWriter('part4')  
writer.add_images('images', grid)  
writer.add_graph(net, [images.cuda()])
```

```
criterion = nn.CrossEntropyLoss()  
# optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.9)  
optimizer = optim.Adam(net.parameters(), lr=0.001, amsgrad=True) # Q3 & Q4
```

```
print('Start training...')  
for epoch in range(MAX_EPOCH): # loop over the dataset multiple times
```

```
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):
```

### CS535 DL Assignment3 Dongkyu Kim

```
# get the inputs
inputs, labels = data

# wrap them in Variable
inputs, labels = Variable(inputs).cuda(), Variable(labels).cuda()

# zero the parameter gradients
optimizer.zero_grad()

# forward + backward + optimize
outputs = net(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

# print statistics
running_loss += loss.item()
if i % 500 == 499: # print every 2000 mini-batches
    print(' Step: %5d avg_batch_loss: %.5f' %
          (i + 1, running_loss / 500))
    running_loss = 0.0
print(' Finish training this EPOCH, start evaluating...')
train_loss, train_acc = eval_net(trainloader)
test_loss, test_acc = eval_net(testloader)
print('EPOCH: %d train_loss: %.5f train_acc: %.5f test_loss: %.5f test_acc %.5f' %
      (epoch+1, train_loss, train_acc, test_loss, test_acc))

#writer.add_scalar('train_loss', train_loss,epoch)
#writer.add_scalar('train_acc', train_acc,epoch)
#writer.add_scalar('test_loss', test_loss,epoch)
#writer.add_scalar('test_acc', test_acc,epoch)

writer.add_scalars('loss', {'train_loss': train_loss,
                           'test_loss': test_loss},epoch)

writer.add_scalars('acc', {'train_acc': train_acc,
                           'test_acc': test_acc},epoch)

# writer.close()
print('Finished Training')
print('Saving model...')
torch.save(net.state_dict(), 'mytraining4.pth')
```